Hadoop Streaming Program using Python

_____MAPPER_____

1> *make a file named mapper.py and paste below python code for mapper in it*

$ nano mapper.py

```
#!/usr/bin/env python

import sys


for line in sys.stdin:

    line = line.strip()

    words = line.split()

    for word in words:

        print '%s\t%s' % (word, 1)
```

--------understanding above code----------

#[for line in sys.stdin:] described that input comes from standard input (STDIN).
Standard input(stdin), is the source of input data for python ,

#[line = line.strip()] removes extra spaces

#[words = line.split()] splits line into words

#[for word in words:] increases counters

#[print '%s\t%s' % (word, 1)] will write the result to (stdout) . This output will
input for reducer

2> *Grant permission to mapper.py*

$ chmod 744 /home/ubuntu/mapper.py

_____REDUCER_____

3> *make a file named reducer.py and paste below python code for reducer in it*

$ nano reducer.py

```
#!/usr/bin/env python

from operator import itemgetter
import sys

current_word = None
current_count = 0
word = None


for line in sys.stdin:

    line = line.strip()


    word, count = line.split('\t', 1)


    try:
        count = int(count)
    except ValueError:

        continue


    if current_word == word:
        current_count += count
    else:
        if current_word:

            print '%s\t%s' % (current_word, current_count)
        current_count = count
        current_word = word


if current_word == word:
    print '%s\t%s' % (current_word, current_count)
```

----understanding above code----

#The code in reducer.py will read results of mapper.py through standard input so , output
of mapper.py and input of reducer.py must match .

#[word, count = line.split('\t', 1)] will parse input got from mapper

#[try:
count = int(count)
except ValueError:] will convert count which is in currently string format to int
because count is going to be a number , i.e int.

#The [continue] statement after the code will ignore the line if count was not the number , i.e int

#[if current_word == word:
current_count += count
else:
if current_word:] here if works because hadoop sorts map output i.e word before it is passed to the reducer

#[print '%s\t%s' % (current_word, current_count)
current_count = count
current_word = word] this will write result to standard output (STDOUT)

4> *Grant all permission to reducer.py*

$ chmod 744 /home/ubuntu/reducer.py

_____RUNNING PYTHON CODE ON HADOOP_____

5> *first copy the files that has to be Processed from our local file system to Hadoop's HD*FS.

$ hadoop fs -put <filename> <input>

6> *run hadoop streaming jar file which will allow python code on hadoop followed by mapper reducer input and output*

$ hadoop jar /usr/local/hadoop/contrib/streaming/hadoop-streaming-1.2.1.jar -file /home/ubuntu/mapper.py -mapper /home/ubuntu/mapper.py -file /home/ubuntu/reducer.py -reducer /home/ubuntu/reducer.py -input in -output out1

----------Understanding above command----------

Here -file takes File/dir to be shipped in the Job jar file -input takes DFS input file for the Map step .
-mapper takes the streaming command to run map steps . -reducer takes the streaming command to run reduce step